

Grididdy: Robot Navigation Using AI Models

1 Perception

The robot uses two sensors:

- **Camera:** Detects the exact positions of adjacent wall tiles deterministically.
- **Magic Sensor:** Returns 1 if at least one adjacent tile is a danger tile. It does not indicate which one, and it cannot detect the goal.

2 State Estimation Using Hidden Markov Model (HMM)

We model the world with an HMM:

- Hidden state X_t : configuration of danger tiles
- Observation E_t : 1 if danger nearby, else 0
- Transition model: $P(X_t | X_{t-1})$ (static environment, so $X_t = X_{t-1}$)
- Sensor model: $P(E_t | X_t)$

We estimate:

$$P(X_t | E_{1:t})$$

to update beliefs about where danger tiles are.

3 State Creation

Each tile (x, y) has a label:

$$m_{x,y} \in \{\text{UNKNOWN, SAFE, WALL, SUSPECTED_DANGER, CONFIRMED_DANGER, GOAL}\}$$

Full robot state:

$$S_t = (x_t, y_t, \{m_{x,y}\}, \{b_{x,y}\})$$

where $b_{x,y}$ is the belief (probability) that tile (x, y) is dangerous.

4 Knowledge Update using Bayesian Inference and Propositional Logic

When the sensor triggers ($E_t = 1$), update belief of adjacent unknowns using Bayes' Rule:

$$P(D_{x,y} | E_t = 1) = \frac{P(E_t = 1 | D_{x,y})P(D_{x,y})}{\sum_j P(E_t = 1 | D_{x_j,y_j})P(D_{x_j,y_j})}$$

If only one adjacent tile is unknown, and sensor is active:

$$m_{x,y} \leftarrow \text{CONFIRMED_DANGER}, \quad b_{x,y} = 1.0$$

5 Reasoning with Bayesian Dynamic Network

The joint probability model over time is:

$$P(X_{0:t}, E_{1:t}) = P(X_0) \prod_{i=1}^t P(X_i | X_{i-1}) P(E_i | X_i)$$

We use filtering to maintain:

$$P(X_t | E_{1:t})$$

6 State Space Search

The grid is treated as a graph. Each tile is a node. We prune:

- WALL tiles
- CONFIRMED_DANGER tiles

Cost function:

$$\text{cost}(x, y) = 1 + \lambda \cdot b_{x,y}$$

where λ is a penalty multiplier (e.g., 100).

7 Action

Actions:

$$A = \{\text{UP, DOWN, LEFT, RIGHT}\}$$

The policy is derived using the Bellman equation:

$$U(S) = R(S) + \gamma \max_a \sum_{S'} P(S' | S, a) U(S')$$

Optimal action:

$$\pi^*(S_t) = \arg \max_a \sum_{S'} P(S' | S_t, a) U(S')$$

Rewards:

- Reaching goal: +1000
- Stepping on danger: -1000
- Any movement: -1

Full Execution Flow (Simplified)

1. Initialize the environment:

- Place 6 wall tiles, 6 danger tiles, 1 goal, and the robot.

2. Perceive surroundings:

- Use camera to mark adjacent WALLs.
- Use magic sensor to detect nearby danger (if any).

3. Update beliefs:

- If sensor triggers:
 - If only 1 adjacent UNKNOWN tile: mark it as `CONFIRMED_DANGER` and set belief = 1.
 - If multiple: mark them as `SUSPECTED_DANGER` and increment their belief by +0.3.
- If no danger detected: downgrade nearby D tiles to `SAFE`.

4. Track visits:

- Count how many times each tile has been visited.

5. Plan path using BFS:

- Avoid WALL and `CONFIRMED_DANGER` tiles.
- Cost of each tile is:

$$\text{cost}(x, y) = b_{x,y} + 0.2 \cdot \text{visit_count}_{x,y}$$

- Choose the path with the lowest total cost.

6. Move the robot:

- Step to the next tile in the best path.
- Mark it `SAFE` and increment visit count.

7. Repeat until goal reached or stuck:

- Stop if the goal is found or no valid moves remain.